
Mathematics of Neural Networks
winter semester 2021/2022
exercise sheet 1

Exercise 1: ((1+1+1)+2+1 points)

- a) Construct feedforward neural networks based on ReLU activations that implement
 - (i) logical AND (\wedge) of two logical variables,
 - (ii) logical OR (\vee) of two logical variables,
 - (iii) logical XOR of two logical variables.Here, $0 \in \mathbb{R}$ encodes **False** and $1 \in \mathbb{R}$ encodes **True**, e.g., $0 \wedge 0 = 0$, $1 \vee 0 = 1$.
- b) Can you find three feedforward neural networks based on ReLU activations that implement these three logical binary operators but with only *one* hidden layer such that the first weight matrix and first bias vector is the *same* for all three?
- c) Prove that there does not exist a feedforward neural network based on ReLU activations that implements XOR with zero hidden layers.

Exercise 2: (2+2 points) Read the Wikipedia entry

https://en.wikipedia.org/wiki/Kolmogorov%E2%80%93Arnold_representation_theorem

and interpret the Kolmogorov–Arnold representation theorem (a) and the variant by Sprecher (b) each as a special type of feedforward neural network.

Exercise 3: (1+1+2 points)

- a) Read the Wikipedia entry on the Basic Linear Algebra Subprograms (BLAS),
https://en.wikipedia.org/wiki/Basic_Linear_Algebra_Subprograms.
- b) Find out which variant of the BLAS your installed variant of NumPy uses by invoking

```
import numpy as np
np.__config__.show()
```
- c) Implement a Python script that computes for given $n, k \in \mathbb{N}$ the product of two matrices $\mathbf{A} \in \mathbb{R}^{n \times n}$ and $\mathbf{B} \in \mathbb{R}^{n \times k}$ using NumPy and
 - (i) nk scalar products (BLAS LEVEL 1),
 - (ii) k matrix-vector products (BLAS LEVEL 2),
 - (iii) one matrix-matrix product (BLAS LEVEL 3).How fast is each variant, say, e.g., for $n = 10.000$ and $k = 100$? You might have to reduce these numbers depending on your computer.

Exercise 4: ((2+2+2)+2+2 points)

- a) Implement a feedforward neural network in Python.

- (i) Implement the **ReLU** activation function as class in the provided **activation.py** script which has a method **evaluate(self, x)** that performs the application of the **ReLU** activation function
 - (ii) Implement a class **DenseLayer** in the given **layers.py** script which
 - is initialized with integers specifying the number of inputs and outputs, and an activation function (which is **ReLU** by default)
 - has the attributes **W** and **b** for the weight matrix and the bias of this layer
 - has the method **evaluate(self, a)** that performs the evaluation on the
 - has the methods **set_weights** and **set_bias** for setting the weight matrix and the bias of a layer.
 - (iii) Implement a class **SequentialNet** in the provided **networks.py** script which
 - has the attributes **layers** that stores all layers of the network and an integer **no** indicating the current number of outputs
 - is initialized by an integer indicating the number of inputs and an (optional) list of layers
 - has the method **evaluate(self,x)** that performs the feed forward with input **x**.
- b) Write a Python script that tests the feed forward process for some given neural net. You may use your results from exercise 1.
- c) (optional) Add a method **draw()** to the class **SequentialNet** that draws the neural network using circles for neurons and lines between them for the connecting weights.