
Mathematics of Neural Networks
winter semester 2021/2022
exercise sheet 3

Exercise 1: (4 points)

- a) Add different weight initialization schemes to `initializers.py` as treated in the lecture:
- (i) A class `RandnAverage` that provides the method `wfun(self, m, n)` that returns an $m \times n$ random matrix with normally distributed entries and variance $2/(m+n)$.
Hint: Use `numpy.random.randn` and multiply the result by the square root of the desired variance.
 - (ii) A class `RandAverage` that provides the method `wfun(self, m, n)` that returns an $m \times n$ random matrix with uniformly distributed entries on the interval $[r, r]$, where the radius is given by $r := \sqrt{6/(m+n)}$.
Hint: Use `numpy.random.rand`. The method `rand` generates random numpy arrays with entries uniformly distributed in the interval $[0, 1]$. Shift and scale the results in an appropriate way.
 - (iii) A class `RandnOrthonormal` that provides the method `wfun(self, m, n)` that returns an $m \times n$ random orthonormal matrix.
Hint: Use `numpy.linalg.svd` or `numpy.linalg.qr` as sketched in the lecture to obtain a random orthonormal matrix from a random matrix.
- b) On initialization of a `DenseLayer` add the attribute `initializer`, which is given as argument on initialization and should be `RandnAverage()` by default.
The weight matrix should be initialized by `wfun` from the initializer and scaled by a factor depending on the activation function. This factor can be obtained from the attribute `factor` of the activation function `afun`.

Exercise 2: (4 points)

- a) Add as many other activation functions to `activations.py` as you like. e.g., those mentioned in the lecture (See lecture notes, section 1.3¹). Use classes and subclasses to group activation functions. See the following listing for an example using the already implemented activation functions:

```
1 import numpy as np
2
3 # Heaviside family activation functions
4 class HeavisideLike:
5
6     def __init__(self):
7
8         self.factor = 1.0
9
10 """
```

¹Version from October 20th, 2021

```

11 TODO Implement the following activation functions from the
12     Heaviside activation function family:
13         - Heaviside function
14         - Modified Heaviside function
15         - Logistic function
16         - Exp function (as mentioned on the exercise sheet)
17 """
18
19 # Sign family activation functions
20 class SignLike:
21
22     def __init__(self):
23
24         self.factor = 1.0
25
26 """
27 TODO Implement the following activation functions from the
28     Sign activation function family:
29         - Sign function
30         - TanH function
31         - SoftSign function
32 """
33 # ReLU family activation functions
34 class ReLULike:
35
36     def __init__(self):
37
38         self.factor = np.sqrt(2)
39
40 class ReLU(ReLULike):
41
42     def __init__(self):
43
44         super().__init__()
45         self.data = None
46         self.name = 'ReLU'
47
48     def evaluate(self, x):
49
50         self.data = x
51         return x.clip(min = 0)
52
53     def backprop(self, delta):
54
55         return (self.data >= 0) * delta
56
57 """
58 TODO Add the following activation functions from the
59     ReLU activation function family:
60         - leaky ReLU function
61         - ELU function
62         - SoftPlus function
63         - Swish function
64 """
65 # Abs family activation functions
66 class AbsLike:

```

```

67
68     def __init__(self):
69
70         self.factor = 1.0
71
72 class Abs(AbsLike):
73
74     def __init__(self, alpha=0.0):
75
76         super().__init__()
77         self.data = None
78         self.name = 'Abs'
79         self.alpha = alpha
80
81     def evaluate(self, x):
82
83         self.data = x
84         if(self.alpha == 0.0):
85             return np.abs(x)
86         else:
87             return np.sqrt(x*x + self.alpha**2) - self.alpha
88
89     def backprop(self, delta):
90
91         if(self.alpha == 0.0):
92             return np.sign(self.data)*delta
93         else:
94             return self.data/np.sqrt(self.data**2 + self.alpha**2) * delta
95
96 """
97 TODO Add the following activation functions from the
98     Abs activation function family:
99     - L0Co function
100     - Twist function
101     - SoftAbs function
102 """

```

Do you have ideas for activation functions different from the ones mentioned in lecture? You could search the internet or take a look at those gathered in Wikipedia.

Hint: The initialization factor for activation functions of one family can be set by calling `super().__init__()`

- b) Expand example 5c) of exercise sheet 2 to use different activation functions. What do you observe? What happens when you use a function that is generally not used as an activation function like $\text{Exp}(x) := e^x$?

Exercise 3: (4 points)

- Add the optimizer Adam (Lecture notes, section 2.5.7²) as class `Adam` to the `optimizers.py` script. It should work similar to the `SGD` class.
- Add momentum to the `SGD` class (See lecture notes, section 2.5.2).

Exercise 4: (4 points) You can add a penalty term for the size of weights by using L^2 regularization to change the cost function. The resulting effect is often called *weight decay*.

²Version from October 20th, 2021

- a) Compare plain SGD using L^2 regularization with SGD using momentum by calculating the first three respective updates of a single weight matrix by hand.
- b) What differences do you observe?
- c) Without directly calculating it, what do you expect for other optimizers like Adam?

Exercise 5: (4 points)

- a) Implement L^2 regularization in our code:
 - (i) Add a class `L2Regularizer` to `layers.py`:
 - It is initialized with a regularization parameter `l2`.
 - It provides a method `update(self, p, dp)` which adds `p` scaled by the regularization parameter `l2` to the derivatives computed with backpropagation.
 - (ii) Add an (optional) attribute `kernel_regularizer` to the `DenseLayer` class which is initialized as `None` by default.
 - (iii) Add the regularization update as additional step in the backpropagation method `backprop` of `DenseLayer`.
- b) Compare the effect of the regularization on plain SGD, SGD with momentum, and Adam by expanding example 5c) of exercise sheet 2. What do you observe?