
Mathematics of Neural Networks
 winter semester 2021/2022
 exercise sheet 5, solutions

Exercise 1: (3 points) In standard implementations, e.g., in TensorFlow/Keras, due to Python's way of storing multiarrays (the last index moves fastest), (activated) minibatches \mathbf{A} of $n \in \mathbb{N}$ images with $c \in \mathbb{N}$ channels of height $h \in \mathbb{N}$ and width $w \in \mathbb{N}$ are stored in the 4D-tensor format

$$\mathbf{A} \in \mathbb{R}^{n \times c \times h \times w},$$

the so-called NCHW format. The filter bank \mathbf{F} of $m \in \mathbb{N}$ filters (m feature maps) of width $f_w \in \mathbb{N}$ and height $f_h \in \mathbb{N}$ that sums over the aforementioned channels c is stored in the 4D-tensor format

$$\mathbf{F} \in \mathbb{R}^{m \times c \times f_h \times f_w}.$$

For every feature map there exists a bias, thus the bias is a vector $\mathbf{b} \in \mathbb{R}^m$.

Derive the backpropagation and kernel update of the feedforward step

$$\mathbf{Z}(i, j, :, :) = b(j)\mathbf{E} + \sum_{k=1}^c \mathbf{A}(i, k, :, :) * \mathbf{F}(j, k, :, :), \quad \mathbf{A}_{\text{new}} = a(\mathbf{Z})$$

for some activation function a by hand. Here \mathbf{E} is a matrix comprising ones.

Hint: Consider backpropagation for one step with given $\Delta \in \mathbb{R}^{n \times m \times z_h \times z_w}$ and $\Delta^{\text{pre}} \in \mathbb{R}^{n \times c \times h \times w}$ to be evaluated. Make use of the multidimensional chain rule. Note, that for the feedforward step above one can show

$$\frac{\partial C}{\partial \mathbf{A}(i, k, :, :)} = \sum_{j=1}^m \frac{\partial C}{\partial \mathbf{Z}(i, j, :, :)} *_{\text{full}} \mathbf{F}(j, k, :, : -1, : -1)$$

for all $i \in \{1, \dots, n\}$, $k \in \{1, \dots, c\}$

Solution: We use the chain rule and partial derivatives. Due to the dimensions we use multi-indices. We define analogously to the backpropagation for dense layers the variables Δ^{pre} of the previous layer, which has to be computed, and Δ of the current layer, which we assume to be given,

$$\Delta^{\text{pre}} \in \mathbb{R}^{n \times c \times h \times w}, \quad \Delta \in \mathbb{R}^{n \times m \times z_h \times z_w},$$

where $z_h := h - f_h + 1$ and $z_w := w - f_w + 1$, by

$$(\Delta^{\text{pre}})_i := \frac{\partial C}{\partial (\mathbf{Z}^{\text{pre}})_i}, \quad (\Delta)_i := \frac{\partial C}{\partial (\mathbf{Z})_i}, \quad \mathbf{i} \in \mathbb{N}^4.$$

Here, $\mathbf{Z}^{\text{pre}} \in \mathbb{R}^{n \times c \times h \times w}$ is the non-activated \mathbf{A} of the previous layer, i.e., either for some activation function a^{pre} we have $\mathbf{A} = a^{\text{pre}}(\mathbf{Z}^{\text{pre}})$, or for the input we have $\mathbf{Z}^{\text{pre}} = \mathbf{X}$, e.g., a^{pre} is the identity.

To derive backpropagation, i.e., compute Δ^{pre} from Δ , we use the chain rule twice, first,

$$(\Delta^{\text{pre}})_i = \frac{\partial C}{\partial \mathbf{Z}_i^{\text{pre}}} = \sum_j \frac{\partial \mathbf{A}_j}{\partial \mathbf{Z}_i^{\text{pre}}} \frac{\partial C}{\partial \mathbf{A}_j} = (a^{\text{pre}})'(\mathbf{Z}_i^{\text{pre}}) \cdot \frac{\partial C}{\partial \mathbf{A}_i}$$

which gives the part $(a^{\text{pre}})'(\mathbf{Z}^{\text{pre}})$ that will be computed in the previous layer. The part that will be computed in the current layer is developed further using the chain rule for the second time,

$$\frac{\partial C}{\partial \mathbf{A}_i} = \sum_j \frac{\partial \mathbf{Z}_j}{\partial \mathbf{A}_i} \frac{\partial C}{\partial \mathbf{Z}_j} = \sum_j \frac{\partial \mathbf{Z}_j}{\partial \mathbf{A}_i} \Delta_j.$$

Now we consider the indices,

$$\frac{\partial C}{\partial \mathbf{A}(i_1, i_2, i_3, i_4)} = \sum_{j_1, j_2, j_3, j_4} \frac{\partial \mathbf{Z}(j_1, j_2, j_3, j_4)}{\partial \mathbf{A}(i_1, i_2, i_3, i_4)} \Delta(j_1, j_2, j_3, j_4).$$

The derivative is zero when $j_1 \neq i_1$, since

$$\mathbf{Z}(j_1, j_2, :, :) = b(j_2)\mathbf{E} + \sum_{i_2=1}^c \mathbf{A}(j_1, i_2, :, :) * \mathbf{F}(j_2, i_2, :, :)$$

does not depend on $\mathbf{A}(i_1, :, :, :)$ for $j_1 \neq i_1$, thus we fix $j_1 = i_1$. The second index implements a sum over all channels, the derivative of a sum is the sum of derivatives, thus we split sums to obtain

$$\frac{\partial C}{\partial \mathbf{A}(i_1, i_2, i_3, i_4)} = \sum_{j_2} \left(\sum_{j_3, j_4} \frac{\partial \mathbf{Z}(i_1, j_2, j_3, j_4)}{\partial \mathbf{A}(i_1, i_2, i_3, i_4)} \Delta(i_1, j_2, j_3, j_4) \right).$$

For fixed indices i_1, i_2, j_2 we have the derivative of the 2D convolution of the matrices $\mathbf{A}(i_1, i_2, :, :)$ with the 2D kernel $\mathbf{F}(j_2, i_2, :, :)$ applied to the matrix $\Delta(i_1, j_2, :, :)$. This is given by the full convolution of the doubly flipped (rotated) kernel with $\Delta(i_1, j_2, :, :)$, i.e., setting for simplicity $i = i_1, j = j_2, k = i_2$,

$$\frac{\partial C}{\partial \mathbf{A}(i, k, :, :)} = \sum_{j=1}^m \Delta(i, j, :, :) *_{\text{full}} \mathbf{F}(j, k, :: -1, :: -1).$$

This quantity is handed to the previous layer.

To compute the bias and kernel update, we first compute the correct Δ from the input of the next layer by componentwise multiplication with $a'(\mathbf{Z})$, compare with the layer-wise approach for dense layers and the above remarks.

The bias is easily obtained, as

$$\mathbf{Z}(i, j, :, :) = b(j)\mathbf{E} + \sum_{k=1}^c \mathbf{A}(i, k, :, :) * \mathbf{F}(j, k, :, :)$$

does only depend in the second position on $b(j)$, thus

$$\frac{\partial C}{\partial b(j)} = \sum_{i, k, \ell} \frac{\partial C}{\partial \mathbf{Z}(i, j, k, \ell)} \frac{\partial \mathbf{Z}(i, j, k, \ell)}{\partial b(j)} = \sum_{i, k, \ell} \Delta(i, j, k, \ell).$$

The computation of all updates of biases $b(j)$ for $j = 1, \dots, m$ can be implemented fast and simultaneously as `np.sum(Delta, axis=(0, 2, 3))` in Python.

The update of the kernel bank is defined by

$$(d\mathbf{F})_i := \frac{\partial C}{\partial \mathbf{F}_i} = \sum_{\mathbf{j}} \frac{\partial \mathbf{Z}_{\mathbf{j}}}{\partial \mathbf{F}_i} \frac{\partial C}{\partial \mathbf{Z}_{\mathbf{j}}} = \sum_{j_1, j_2, j_3, j_4} \frac{\partial \mathbf{Z}(j_1, j_2, j_3, j_4)}{\partial \mathbf{F}(i_1, i_2, i_3, i_4)} \Delta(j_1, j_2, j_3, j_4).$$

As $\mathbf{Z}(:, j_2, :, :)$ does not depend on $\mathbf{F}(i_1, :, :, :)$ for $i_1 \neq j_2$ we set $i_1 = j_2$ and split the remaining sums,

$$d\mathbf{F}(j_2, i_2, i_3, i_4) = \sum_{j_1} \left(\sum_{j_3, j_4} \frac{\partial \mathbf{Z}(j_1, j_2, j_3, j_4)}{\partial \mathbf{F}(j_2, i_2, i_3, i_4)} \Delta(j_1, j_2, j_3, j_4) \right).$$

For simplicity, we set $j = j_2$, $k = i_2$, and $i = j_1$,

$$d\mathbf{F}(j, k, i_3, i_4) = \sum_i \left(\sum_{j_3, j_4} \frac{\partial \mathbf{Z}(i, j, j_3, j_4)}{\partial \mathbf{F}(j, k, i_3, i_4)} \Delta(i, j, j_3, j_4) \right).$$

For fixed i, j, k the term in parentheses corresponds to the (known) derivative of the convolution of $\mathbf{A}(i, k, :, :)$ with $\mathbf{F}(j, k, :, :)$ applied to $\Delta(i, j, :, :)$, i.e.,

$$d\mathbf{F}(j, k, :: -1, :: -1) = \sum_i \mathbf{A}(i, k, :, :) * \Delta(i, j, :: -1, :: -1).$$

This finishes the solution.

Exercise 2: (5 points) Add a class `Conv2DLayer()` that is initialized by

`Conv2DLayer(tensor, fshape, afun=ReLU(), optim=SGD()),`

to `layers.py` where

- `tensor` is a tuple (c, h, w) of channels, height, and width of the (activated) images on input, e.g., for MNIST $(1, 28, 28)$ and
- `fshape` is a tuple (m, f_h, f_w) describing the filter bank of m filters of size $f_h \times f_w$, e.g., for a typical filter bank $(10, 3, 3)$.

The class should store additionally

- `afun`: the activation function a , default value `ReLU()`,
- `optim`: the optimizer used on this layer, default value `SGD()`,
- `initializer`: the initializer for the filter bank with `RandnAverage()` as default value
- `f`: the filter bank initialized with the method `ffun` from the initializer and scaled by `self.afun.factor`,
- `b`: the bias, zero on initialization,
- `__a`: the input \mathbf{a} from the last evaluation, initialized by `None`,
- `__z`: the result $\mathbf{z} = \mathbf{a} * \mathbf{f} + \mathbf{b}$ after applying the filter bank, initialized by `None`,
- `df`: the update of `f` from the backpropagation, zero on initialization,
- `db`: the update of `b` from the backpropagation, zero on initialization.

Add the following methods, using `convolve2d` of the SciPy package wherever necessary.

- `evaluate(self, a)`: Evaluate the feedforward step (see lecture notes, p.106¹, or exercise 1)

¹Version from 01.11.2021

- `update(self)`: updates the filter bank and the bias using the method `update()` of the optimizer.

to the new class. A method `backprop` is already implemented in `layers.py`. Check your implementation using gradient checking which is provided to you in the file `layers.py`.

Solution: A possible implementation can be found in the following listing:

```

1 class Conv2DLayer:
2
3     def __init__(self, tensor, fshape, afun=None, optim=None, initializer=
4         None):
5
6         if afun is None:
7             self.afun = ReLU()
8         else:
9             self.afun = afun
10
11         if optim is None:
12             self.optim = SGD()
13         else:
14             self.optim = optim
15
16         if initializer is None:
17             self.initializer = RandnAverage()
18         else:
19             self.initializer = initializer
20
21         self.tensor = tensor
22         self.fshape = fshape
23
24         m, fh, fw = fshape
25
26         c, h, w = tensor
27         self.f = self.initializer.ffun(m, c, fh, fw)
28         self.f *= self.afun.factor
29         self.b = np.zeros(m)
30
31         self.__a = None
32         self.__z = None
33
34         self.df = np.zeros_like(self.f)
35         self.db = np.zeros_like(self.b)
36
37     def evaluate(self, a):
38
39         n, c, h, w = a.shape
40         m, fh, fw = self.fshape
41
42         self.__a = a
43
44         self.__z = np.zeros((n, m, h - fh + 1, w - fw + 1))
45         for j in range(m):
46             for i in range(n):
47                 self.__z[i, j, :, :] += self.b[j]
48                 for k in range(c):
49                     self.__z[i, j, :, :] += convolve2d(self.__a[i, k, :, :], \

```

```

50                                     mode='valid')
51
52     return self.afun.evaluate(self.__z)
53
54     def backprop(self, delta):
55
56         n, c, h, w = self.__a.shape
57         m, fh, fw = self.fshape
58         # Compute a'(z)*delta
59         delta = self.afun.backprop(delta)
60
61         # Compute bias change
62         self.db = np.sum(delta, axis=(0,2,3))
63
64         # Compute df = delta * f~. df has shape (m, c, fh, fw)
65         for k in range(c):
66             for j in range(m):
67                 for i in range(n):
68                     self.df[j,k,:,:] += convolve2d(self.__a[i,k,:,:-1,:,:-1],
69                                                     delta[i,j,:,:],
70                                                     mode = 'valid')
71
72         # Update delta
73         # The new delta has the same shape as __a: (n,c,h,w)
74         delta2 = np.zeros_like(self.__a)
75         for k in range(c):
76             for i in range(n):
77                 for j in range(m):
78
79                     delta2[i, k,:,:] += convolve2d(delta[i,j,:,:],
80                                                     self.f[j,k,:,:-1,:,:-1])
81
82     return delta2
83
84     def update(self):
85
86         self.optim.update([self.f, self.b],
87                           [self.df, self.db])

```

Exercise 3: (3 points) Add the classes `Pool2DLayer()` and `FlattenLayer()` to `layers.py`.

a) `Pool2DLayer()` should be initialized by `Pool2DLayer(area)`, where `area` is a tuple (h, w) describing the window where to compute the maximum, e.g., $(2, 2)$, and store the following additional variables

- **shape:** a tuple n, c, h, w of number of images, channels, height, and width of the (activated) images on input, initialized by `None`,
- **mask:** a tensor of boolean values storing the information which indices have “won” the max pooling.

Add the methods

- **evaluate(self, a)**, store the dimensions of the input `a` in `shape`, reshape your input as necessary, save the information which indices have “won” in `mask` and return the result of the max pooling,
- **backprop(self, delta)**, use the information in `mask` and suitable reshaping to return the correct `delta`.

Also add the method `add_pool2D(self, area, strict=False)` to our `SequentialNet` class.

- b) `FlattenLayer()` should be initialized by `FlattenLayer()`. This layer should map input tensors of size (n, c, h, w) to a matrix of size $(n, c * h * w)$ to be consistent with the storage layout in `DenseLayer()`. It should only store
- **shape**: a tuple n, c, h, w of number of images, channels, height, and width of the (activated) images on input, initialized by `None`.

Add the methods

- `evaluate(self, a)`, store the dimensions of the input `a` in `shape` and return a matrix consistent with the storage layout in `DenseLayer()`,
- `backprop(self, delta)`, returns a suitable reshaped delta.

Also add the method `add_flatten(self)` to our `SequentialNet` class.

- c) Add the method

- `update(self)`.

to both classes. What does it do?

- d) What happens in your backpropagation for the class `Pool2DLayer()`, when the evaluation before has identified more than one maximum in one area? Add the boolean parameter `strict` to your class `Pool2DLayer()`, that is either `True` or `False` with default value `False`. If it is true the backpropagation distributes the corresponding delta evenly over every index where the evaluation of the pooling layer had identified a maximum before by dividing the input delta by the number of maximums. If it is false, it does the same but without dividing it by the number of maximums.

Check your implementation by using the code already provided in `layers.py`.

Solution: A possible implementation can be found in the following listing:

```

1 class Pool2DLayer:
2
3     def __init__(self, area, strict=False):
4
5         self.area = area
6         self.shape = None
7         self.mask = None
8         self.strict = strict
9
10    def evaluate(self, a):
11
12        n, c, h, w = self.shape = a.shape
13
14        ph, pw = self.area
15        nh, nw = h // ph, w // pw
16        oh, ow = h % ph, w % pw
17
18        # Reduce a, if height and width are odd
19        a_reduced = a[:, :, h-oh, :w-ow]
20
21        # Reshape reduced a to find maxima
22        a_reshaped = a_reduced.reshape(n, c, nh, ph, nw, pw)
23
24        # Find maxima and create mask
25        z = a_reshaped.max(axis=(3,5))
26        z_newaxis = z[:, :, :, np.newaxis, :, np.newaxis]
27        self.mask = (a_reshaped == z_newaxis)

```

```

28         return z
29
30     def backprop(self, delta):
31
32         n, c, h, w = self.shape
33         ph, pw      = self.area
34
35         # New delta should be of shape (n, c, h, w).
36         # It can be recovered from the reshaped form
37         da_reshaped = np.zeros((n, c, h//ph, ph, w//pw, pw))
38
39         # Broadcast delta to shape of da_reshaped
40         delta_newaxis = delta[:, :, :, np.newaxis, :, np.newaxis]
41         delta_broadcast, _ = np.broadcast_arrays(delta_newaxis, da_reshaped)
42
43         # Set points, where maximums were found to 1
44         da_reshaped[self.mask] = delta_broadcast[self.mask]
45
46         # Strict mode for correction of the subgradients in case of multiple
47         # maxima
48         if self.strict:
49             da_reshaped /= np.sum(self.mask, axis=(3,5), keepdims=True)
50
51         # Build da from da_reshaped
52         da = np.zeros(self.shape)
53         dah, daw = h - h%ph, w - w%pw
54         da[:, :, :, dah, :, daw] = da_reshaped.reshape(n, c, dah, daw)
55
56         return da
57
58     def update(self):
59
60         """
61         No update is done for Pooling layers. To keep the update method
62         from
63         SequentialNet, it is included here, but nothing is done
64         """
65         pass
66
67 class FlattenLayer:
68
69     def __init__(self):
70
71         self.shape = None
72
73     def evaluate(self, a):
74
75         self.shape = a.shape
76         # The -1 allows numpy to figure out the second dimension automatically
77         return a.reshape(self.shape[0], -1)
78
79     def backprop(self, delta):
80
81         # Reshape delta from shape (n, c*h*w) to shape (n, c, h, w)
82         return delta.reshape(self.shape)

```

```

82
83     def update(self):
84
85         pass

```

Exercise 4: (3 points) Develop a CNN for the MNIST dataset, using all layers implemented in the exercises above. You can test your network for example using the following skeleton and either our implementation or TensorFlow.

```

1  import numpy          as np
2  import matplotlib.pyplot as plt
3
4  from random           import randrange
5  # If you use TF:
6  # import tensorflow.keras as tfk
7
8  from networks         import SequentialNet
9  from layers           import *
10 from optimizers       import *
11 from activations      import *
12
13 DATA = np.load('mnist.npz')
14 x_train, y_train = DATA['x_train'], DATA['y_train']
15 x_test, y_test   = DATA['x_test'], DATA['y_test']
16 x_train, x_test = x_train / 255.0, x_test / 255.0
17
18 """
19 TODO Implement the network you have developed for exercise 4
20     Note that x_train and x_test are of shape (60000,28,28)
21     and (10000,28,28). You need to add an additional axis.
22     This can be done with np.newaxis, e.g
23     x_test = x_test[:,np.newaxis,:,:]
24 """
25 netz=None
26
27 y_tilde = netz.evaluate(x_test)
28 guess   = np.argmax(y_tilde, 1).T
29 print('accuracy =', np.sum(guess == y_test)/100)
30
31 for i in range(4):
32
33     k = randrange(y_test.size)
34     plt.title('Label is {lb}, guess is {gs}'.format(lb=y_test[k], gs=guess[k]
35     ))
36     plt.imshow(x_test[k], cmap='gray')
37     plt.show()

```

Solution: A possible solution with our implementation can be found in the following listing:

```

1  import numpy          as np
2  import matplotlib.pyplot as plt
3
4  from random           import randrange
5
6  from networks         import SequentialNet
7  from layers           import *

```



```

8 from optimizers import *
9 from activations import *
10
11 DATA = np.load('mnist.npz')
12 x_train, y_train = DATA['x_train'], DATA['y_train']
13 x_test, y_test = DATA['x_test'], DATA['y_test']
14 x_train, x_test = x_train / 255.0, x_test / 255.0
15
16 bs, ep, eta = 1000, 10, .001
17
18 x = x_train[:,np.newaxis,:,:)
19 I = np.eye(10)
20 y = I[y_train,:]
21
22
23 net = SequentialNet((1, 28,28))
24 net.add_conv2D((32,3,3), afun=ReLU(), optim=Adam(eta))
25 net.add_pool2D((2,2))
26 net.add_flatten()
27 net.add_dense(100, afun=ReLU(), optim=Adam(eta))
28 net.add_dense(10, afun=SoftMax(), optim=Adam(eta))
29
30 net.train(x, y, bs, ep)
31
32 y_tilde = net.evaluate(x_test[:,np.newaxis,:,:)
33 guess = np.argmax(y_tilde, 1).T
34 print('accuracy =', np.sum(guess == y_test)/100)
35
36 for i in range(4):
37
38     k = randrange(y_test.size)
39     plt.title('Label is {lb}, guess is {gs}'.format(lb=y_test[k], gs=guess[k]
40     ))
41     plt.imshow(x_test[k], cmap='gray')
42     plt.show()

```

A solution using TensorFlow can be found in the following listing.

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 import tensorflow.keras as tfk
4
5 from random import randrange
6
7 from networks import SequentialNet
8 from layers import *
9 from optimizers import *
10 from activations import *
11
12 DATA = np.load('mnist.npz')
13 x_train, y_train = DATA['x_train'], DATA['y_train']
14 x_test, y_test = DATA['x_test'], DATA['y_test']
15 x_train, x_test = x_train / 255.0, x_test / 255.0
16
17 bs, ep, eta = 1000, 10, .001
18
19 x = x_train[:, :, :, np.newaxis]

```

```

20 I = np.eye(10)
21 y = I[y_train,:]
22
23 layers = [tfk.layers.Conv2D(32, (3, 3),
24                             input_shape=(28,28,1),
25                             activation='relu'),
26           tfk.layers.MaxPool2D((2,2)),
27           tfk.layers.Flatten(),
28           tfk.layers.Dense(100,
29                             activation='relu'),
30           tfk.layers.Dense(10,
31                             activation='softmax')]
32
33 net = tfk.Sequential(layers)
34
35 opt = tfk.optimizers.Adam(learning_rate=eta)
36 net.compile(optimizer=opt,
37             loss='categorical_crossentropy',
38             metrics=['accuracy'])
39
40 net.fit(x, y, batch_size=bs, epochs=ep)
41
42 y_tilde = net.predict(x_test[:, :, :, np.newaxis])
43 guess    = np.argmax(y_tilde, 1).T
44 print('accuracy =', np.sum(guess == y_test)/100)
45
46 for i in range(4):
47
48     k = randrange(y_test.size)
49     plt.title('Label is {lb}, guess is {gs}'.format(lb=y_test[k], gs=guess[k]
50                                                     )))
50     plt.imshow(x_test[k], cmap='gray')
51     plt.show()

```