
Mathematics of Neural Networks
winter semester 2021/2022
exercise sheet 6, solutions

Exercise 1: (4 points) Derive an FFT for $n = 3^k$.

Hint: Similar to the case $n = 2 \cdot m$ in the lecture notes on page 112¹ consider $n = 3 \cdot m$ and decompose the DFT of size n into three DFT's of size m .

Solution: According to the lecture, a DFT for $n = 3 \cdot m$ can be decomposed into three DFTs of size m . Let

$$\omega_\ell := e^{-\frac{2\pi i}{\ell}}, \quad \ell \in \mathbb{N}_0, \quad \omega_\ell^\ell = 1.$$

The DFT is given for $n = 3m$ by

$$\hat{x}_k = \sum_{j=0}^{3m-1} (\omega_{3m}^k)^j x_j, \quad k = 0, \dots, 3m-1.$$

We group the indices j in the DFT into the three distinct classes

$$j = 0 \pmod{3}, \quad j = 1 \pmod{3}, \quad j = 2 \pmod{3},$$

and the indices k into m distinct classes, $k = q \pmod{m}$. Then for $j = 3r + \ell$ with $j = \ell \pmod{3}$ and $k = pm + q$ with $k = q \pmod{m}$,

$$(\omega_{3m}^k)^j = \omega_{3m}^{(pm+q)(3r+\ell)} = \omega_{3m}^{(3mpr+mp\ell+3qr+q\ell)} = \omega_3^{p\ell} \cdot \omega_m^{qr} \cdot \omega_n^{q\ell}.$$

The DFT can be expressed as

$$\begin{aligned} \hat{x}_k &= \sum_{j=0}^{3m-1} (\omega_{3m}^k)^j x_j = \sum_{\ell=0}^2 \omega_3^{p\ell} \cdot \omega_n^{q\ell} \sum_{r=0}^{m-1} \omega_m^{qr} \cdot x_{3r+\ell} \\ &= \begin{cases} \sum_{r=0}^{m-1} (\omega_m^q)^r x_{3r} + \omega_3^0 \cdot \omega_n^q \sum_{r=0}^{m-1} (\omega_m^q)^r x_{3r+1} + \omega_3^0 \cdot \omega_n^{2q} \sum_{r=0}^{m-1} (\omega_m^q)^r x_{3r+2}, & p = 0, \\ \sum_{r=0}^{m-1} (\omega_m^q)^r x_{3r} + \omega_3^1 \cdot \omega_n^q \sum_{r=0}^{m-1} (\omega_m^q)^r x_{3r+1} + \omega_3^2 \cdot \omega_n^{2q} \sum_{r=0}^{m-1} (\omega_m^q)^r x_{3r+2}, & p = 1, \\ \sum_{r=0}^{m-1} (\omega_m^q)^r x_{3r} + \omega_3^2 \cdot \omega_n^q \sum_{r=0}^{m-1} (\omega_m^q)^r x_{3r+1} + \omega_3^4 \cdot \omega_n^{2q} \sum_{r=0}^{m-1} (\omega_m^q)^r x_{3r+2}, & p = 2. \end{cases} \end{aligned}$$

With

$$\omega_3^0 = 1, \quad \omega_3^1 = \omega_3^4 = \omega_3, \quad \omega_3^2 = \overline{\omega_3}$$

¹Version from 01.11.21

we see that

$$\mathbf{F}_{3m}[:, 0, 3, 6, \dots, 1, 4, 7, \dots, 2, 5, 8, \dots] = \begin{pmatrix} \mathbf{I}_m & \boldsymbol{\Omega}_m & \boldsymbol{\Omega}_m^2 \\ \mathbf{I}_m & \omega_3 \boldsymbol{\Omega}_m & \overline{\omega_3} \boldsymbol{\Omega}_m^2 \\ \mathbf{I}_m & \overline{\omega_3} \boldsymbol{\Omega}_m & \omega_3 \boldsymbol{\Omega}_m^2 \end{pmatrix} \begin{pmatrix} \mathbf{F}_m & \mathbf{O} & \mathbf{O} \\ \mathbf{O} & \mathbf{F}_m & \mathbf{O} \\ \mathbf{O} & \mathbf{O} & \mathbf{F}_m \end{pmatrix},$$

where

$$\boldsymbol{\Omega}_m := \text{diag}(\omega_n^0, \omega_n^1, \dots, \omega_n^{m-1}).$$

A recursive implementation is given in the following listing.

```

1 import numpy as np
2 from scipy.linalg import dft
3
4 def fft3(x):
5     n = len(x)
6     if n == 1:
7         y = x
8     else:
9         y = np.zeros(n, dtype='complex')
10        m = n//3
11        omega = np.exp(-2*np.pi*1j/n)
12        d = omega**np.arange(m)
13        om3 = np.exp(-2*np.pi*1j/3)
14        oc3 = om3.conj()
15        z_top = fft3(x[0:n:3])
16        z_mid = d * fft3(x[1:n:3])
17        z_bot = d * d * fft3(x[2:n:3])
18        y[0*m:1*m] = z_top + z_mid + z_bot
19        y[1*m:2*m] = z_top + om3*z_mid + oc3*z_bot
20        y[2*m:3*m] = z_top + oc3*z_mid + om3*z_bot
21    return y
22
23 k = 3
24 x = np.random.randint(0, 10, 3**k)
25 y_fft3 = fft3(x)
26 print("Our variant =\n", y_fft3)
27 y = np.fft.fft(x)
28 print("NumPy's variant =\n", y)
29 y_dft = dft(3**k) @ x
30 print("Multiplication by SciPy's DFT matrix =\n", y_dft)
31 print("Difference between Numpy's and our variant =",
32       np.linalg.norm(y-y_fft3))

```

Exercise 2: (4 points)

- Count operations in the following two algorithms to compute the full convolution $\mathbf{y} = \mathbf{x} * \mathbf{f}$ of $\mathbf{x} \in \mathbb{R}^n$ and $\mathbf{f} \in \mathbb{R}^k$:
 - directly compute the convolution,
 - enlarge the dimension of \mathbf{x} and \mathbf{f} to the smallest $2^\ell \geq n+k-1$ by appending trailing zeros and use the radix-2 FFT/IFFT (w/o the permutation).
- Write a Python script that plots for given $n \in \mathbb{N}$ the operation count of both approaches. For which k is FFT better?

Hint: You can use the following Python code for part (ii) of a).

```

1 import numpy as np
2 from scipy.linalg import dft
3
4 def fft2(x):
5     n = len(x)
6     if n == 1:
7         y = x
8     else:
9         y = np.zeros(n, dtype='complex')
10        m = n//2
11        omega = np.exp(-2*np.pi*1j/n)
12        d = omega**np.arange(m)
13        z_top = fft2(x[0:n:2])
14        z_bot = d * fft2(x[1:n:2])
15        y[0:m] = z_top + z_bot
16        y[m:n] = z_top - z_bot
17    return y

```

Solution: We suppose that $n \geq k$, otherwise we switch the inputs.

- a) Computing the convolution directly involves k multiplications and $k - 1$ additions for every of the $n - k + 1$ outputs that completely overlap. A partial overlap occurs $2(k - 1)$ times, the number of multiplications and additions range between $k - 1$ and 1 and $k - 2$ and zero, respectively, i.e,

$$M(n, k) = k(n - k + 1) + 2 \sum_{j=1}^{k-1} j = nk - k(k - 1) + k(k - 1) = nk,$$

$$A(n, k) = (k - 1)(n - k + 1) + 2 \sum_{j=1}^{k-2} j = nk - (n + k - 1) = (n - 1)(k - 1).$$

where M denotes the number of multiplications and A the number of additions. Counting one addition and one multiplication as one operation, we end up with $\mathcal{O}(nk)$ operations.

- b) We count the operations in the recursive implementation of the following listing:

```

1 import numpy as np
2 from scipy.linalg import dft
3
4 def fft2(x):
5     n = len(x)
6     if n == 1:
7         y = x
8     else:
9         y = np.zeros(n, dtype='complex')
10        m = n//2
11        omega = np.exp(-2*np.pi*1j/n)
12        d = omega**np.arange(m)
13        z_top = fft2(x[0:n:2])
14        z_bot = d * fft2(x[1:n:2])
15        y[0:m] = z_top + z_bot
16        y[m:n] = z_top - z_bot
17    return y
18
19 k = 5

```

```

20 x = np.random.binomial(10, .5, 2**k)
21 y_fft2 = fft2(x)
22 print("Our variant=\n", y_fft2)
23 y = np.fft.fft(x)
24 print("NumPy's variant=\n", y)
25 y_dft = dft(2**k) @ x
26 print("Multiplication by SciPy's DFT matrix=\n", y_dft)
27 print("Difference between Numpy's and our variant=",
28       np.linalg.norm(y-y_fft2))

```

We call `fft2` with a vector of length $m = 2^\ell$. We do not count the integer division nor the computation of ω . We count the computation of the vector of powers of ω as $m/2$ multiplications, together with the $m/2$ elementwise multiplications with these we have m multiplications per call to `fft2`. The computation of the upper and lower half of \mathbf{y} consists of m additions. We call `fft2` twice with a vector of size $m/2$, fourth with a vector of size $m/4$, i.e.,

$$M(m) = m + 2m/2 + 4m/4 + \dots + 2^{\ell-1}m/2^{\ell-1} = \ell 2^\ell,$$

$$A(m) = M(m).$$

As we map every vector that has more elements than $2^{\ell-1}$ to an extended version of size 2^ℓ , the function that counts the operations is piecewise linear. We call `fft2` twice, once for $\mathbf{x} \in \mathbb{R}^n$, once for $\mathbf{f} \in \mathbb{R}^k$. The componentwise multiplication adds 2^ℓ multiplications. We count the single call of `ifft2` as one call of `fft2`, thus we have in total $(3\ell + 1)2^\ell$ multiplications and $3\ell 2^\ell$ additions, which we count together as $(6\ell + 1)2^\ell$ operations.

The two functions are thus given by

$$f_1(n, k) = nk + (n-1)(k-1), \quad f_2(n, k) = (6\ell + 1)\ell 2^\ell, \quad \ell = \lceil \log(n+k-1) \rceil.$$

The following listing gives the plots for selected n and all $k = 1, \dots, n$.

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 for n in range(2,100):
5     k = np.arange(1,n+1)
6
7     direct = n*k + (n-1)*(k-1)
8     ell = np.ceil(np.log2(n+k-1))
9     fft = (6*ell+1)*(2**ell)
10
11     print(direct.shape)
12
13     plt.plot(k,direct,'r-',k,fft,'b.')
14     plt.legend(['direct_convolution', 'convolution_via_fft'])
15     plt.title("n={n}".format(n=n))
16     if n == 80:
17         plt.savefig('theory80.pdf')
18     plt.show()
19     plt.close()

```

For an example, see [Figure 1](#).

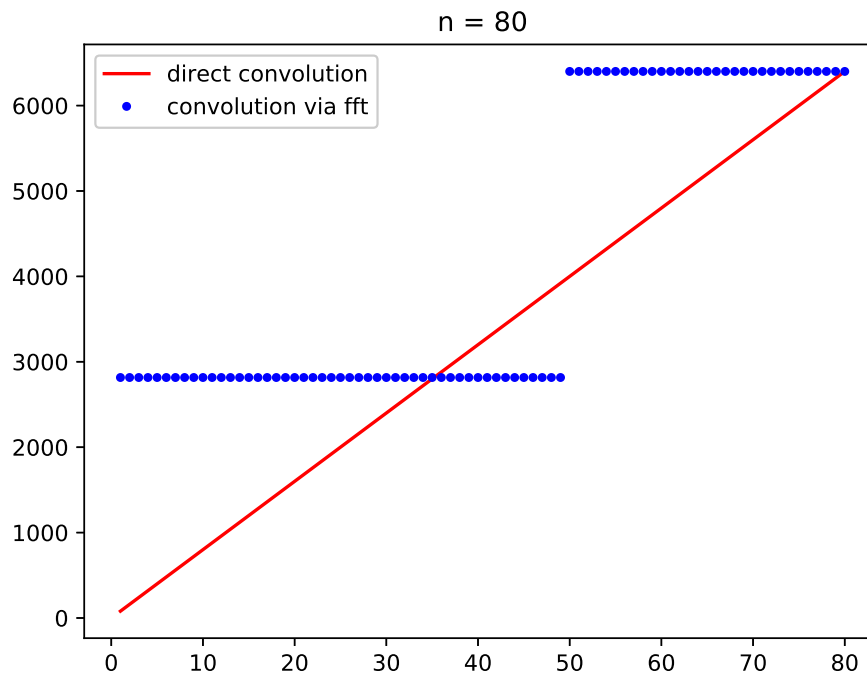


Figure 1: Theoretical comparison of direct and FFT-based convolution, $n = 80$.

Exercise 3: (4 points) In exercise 2 on the previous exercise sheet you implemented evaluation and backpropagation in a convolutional layer using the function `convolve2D()` of the SciPy package. Now we want to utilize the FFT approach and the `im2col` approach to calculate the convolution.

- Add a method `evaluate_fft(self, a)` that implements the evaluation of a 2D convolutional layer with FFT and IFFT. Use the functions `rfft2` and `irfft2` from `scipy.fft`²
- Add a method `evaluate_im2col(self, a)` that implements the evaluation of a 2D convolutional layer with `im2col`. Use the functions `im2col` from the script `utils.py` to map the input to the corresponding `im2col` matrix. Use the attribute `cache` of `Conv2DLayer` to save the `im2col` matrix and the reshaped filterbank. Don't forget to reshape the result before applying the activation function.
- Test your implementation with the code provided in `layers.py`

Solution: A possible solution for the evaluation with FFT and IFFT can be found in the following listing.

```

1  def evaluate_fft(self, a):
2
3      self.__a = a
4
5      n, c, h, w = a.shape
6      m, fh, fw = self.fshape
7      dh, dw = fh - 1, fw - 1
8      zh, zw = h - dh, w - dw
9

```

²More information can be found at <https://docs.scipy.org/doc/scipy/reference/tutorial/fft.html>

```

10     # Apply FFT
11     a_hat = rfft2(a, (h + dh, w + dw))
12     f_hat = rfft2(self.f, (h + dh, w + dw))
13
14     self.__z = np.zeros((n, m, zh, zw))
15     for i in range(n):
16         for j in range(m):
17
18             """
19             Compute the full convolution with applying
20             the componentwise product of a_hat and f_hat.
21             z_hat then has the shape (c, h+dh, w+dw) and has
22             to be restricted and summed over all channels.
23             """
24             z_hat = irfft2(a_hat[i,:,:,:] * \
25                             f_hat[j,:,:,:],
26                             (h + dh, w + dw))
27             self.__z[i,j,:,:] = self.b[j] + \
28                 z_hat[:, dh:-dh, dw:-dw].sum(axis=0)
29
30     return self.afun.evaluate(self.__z)

```

The evaluation with `im2col` is shown in the following listing.

```

1     def evaluate_im2col(self, a):
2
3         self.__a = a
4
5         n, c, h, w = a.shape
6         m, fh, fw = self.fshape
7         zh, zw = h - fh + 1, w - fw + 1
8
9         # Create im2col matrix
10        a_col = im2col(a, fh, fw)
11        # reshape filterbank
12        f_row = self.f[:, :, ::-1, ::-1].reshape(m, -1)
13        self.cache = a_col, f_row
14
15        # BLAS Level 3: GEMM
16        z_mat = f_row @ a_col + self.b.reshape(-1, 1)
17
18        # reshape
19        self.__z = z_mat.reshape(m, n, zh, zw).transpose(1, 0, 2, 3)
20        return self.afun.evaluate(self.__z)

```

Exercise 4: (4 points) Compare our implementation of a neural network with TensorFlow.

- Develop a convolutional neural network for the Fashion MNIST dataset. Try to achieve at least 80 percent accuracy.
- Implement the network with TensorFlow and our library.
- Compare the time needed for **one** epoch of training in both cases. Use `im2col` for evaluation in our implementation. The time for the TensorFlow network is printed while the network is trained.

You may use the following skeleton provided in `skeleton.py`.

```

1 import numpy as np

```

```

2 import matplotlib.pyplot as plt
3 import tensorflow.keras as tfk
4
5 from random import randrange
6 from time import time # For time measuring
7
8 from networks import SequentialNet
9 from layers import *
10 from optimizers import *
11 from activations import *
12
13 DATA = np.load('fashion_mnist.npz')
14 x_train, y_train = DATA['x_train'].reshape(60000,28,28), DATA['y_train']
15 x_test, y_test = DATA['x_test'].reshape(10000,28,28), DATA['y_test']
16 x_train, x_test = x_train / 255.0, x_test / 255.0
17
18 x = x_train[:,np.newaxis,:,:)
19 x_TF = x_train[:,:,:,:np.newaxis]
20
21 bs, ep, eta = 128, 10, .001
22 """
23 Categories for Fashion MNIST. Category i is ct[i].
24 """
25 ct = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress',
26       'Coat', 'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle_Boot']
27
28 """
29 TODO Set up the network with our library
30 """
31 net = SequentialNet((1,28,28))
32
33 """
34 Add the layers to your network. As first hidden layer
35 you can use for example:
36
37 net.add_conv2D((32,3,3),
38               afun=ReLU(),
39               optim=Adam(),
40               initializer=HeUniform(),
41               eval_method='im2col')
42 """
43
44 """
45 The Last layer should be a SoftMax Layer with 10 neurons
46 """
47 net.add_dense(10, afun=SoftMax(),
48               optim=Adam(),
49               initializer=HeUniform())
50 """
51 TODO Set up the network with TensorFlow
52 """
53 input_shape = (28,28,1)
54 net_TF = tfk.Sequential()
55
56 net_TF.add(tfk.Input(shape=input_shape))
57 """

```

```

58 Add the same layers as above to the network.
59 If you used a convolutional layer with 32 3x3 filters
60 as first layer, you can add it with
61 net_TF.add(tfk.layers.Conv2D(32, (3,3),
62                             activation='relu',
63                             kernel_initializer='he_uniform'))
64 """
65
66 net_TF.add(tfk.Dense(10, activation='softmax',
67                     kernel_initializer='he_uniform'))
68
69 opt = tfk.optimizers.Adam(eta)
70 net_TF.compile(optimizer=opt,
71               loss='categorical_crossentropy',
72               metrics=['accuracy'])
73
74
75 start = time()
76 net.train(x, y_train, batch_size=bs, epochs=1)
77 t_train = time() - start
78
79 """
80 TODO Train the TensorFlow network. With metrics=['accuracy'] you
81 get the time needed for training one epoch.
82 """
83
84 y_test = np.argmax(y_test, 1).T
85
86
87
88 y_tilde_TF = net_TF.predict(x_test.reshape(10000, 28, 28, 1))
89 guess_TF    = np.argmax(y_tilde_TF, 1).T
90 print('Accuracy with TensorFlow=', np.sum(guess_TF == y_test)/100)

```

Solution: A solution can be found in the following listing.

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 import tensorflow.keras as tfk
4
5 from random import randrange
6 from time import time # For time measuring
7
8 from networks import SequentialNet
9 from layers import *
10 from optimizers import *
11 from activations import *
12 from initializers import *
13
14 DATA = np.load('fashion_mnist.npz')
15 x_train, y_train = DATA['x_train'].reshape(60000, 28, 28), DATA['y_train']
16 x_test, y_test = DATA['x_test'].reshape(10000, 28, 28), DATA['y_test']
17 x_train, x_test = x_train / 255.0, x_test / 255.0
18
19 x = x_train[:, np.newaxis, :, :]
20 x_TF = x_train[:, :, :, np.newaxis]
21

```



```

22 """
23 Categories for Fashion MNIST. Category i is ct[i].
24 """
25 ct = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress',
26       'Coat', 'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle_Boot']
27
28 """
29 Set up the network with our library
30 """
31 bs, ep, eta = 128, 10, .001
32
33
34 net = SequentialNet((1,28,28))
35 net.add_conv2D((32,3,3),
36               afun=ReLU(),
37               optim=Adam(eta),
38               initializer=
39               HeUniform(),
40               eval_method='im2col')
41 net.add_pool2D((2,2))
42 net.add_conv2D((64,3,3),
43               afun=ReLU(),
44               optim=Adam(eta),
45               initializer=HeUniform(),
46               eval_method='im2col')
47 net.add_pool2D((2,2))
48 net.add_flatten()
49 net.add_dense(2048,
50               afun=ReLU(),
51               optim=Adam(eta),
52               initializer=HeUniform())
53 net.add_dense(10,
54               afun=SoftMax(),
55               optim=Adam(eta),
56               initializer=HeUniform())
57
58
59 """
60 Set up the network with TensorFlow
61 """
62 activation = 'relu'
63 input_shape=(28,28,1)
64
65 net_TF = tfk.Sequential()
66
67 net_TF.add(tfk.Input(shape=input_shape))
68 net_TF.add(tfk.layers.Conv2D(32, (3,3),
69                               activation=activation,
70                               kernel_initializer='he_uniform'))
71 net_TF.add(tfk.layers.MaxPool2D(pool_size=(2,2)))
72 net_TF.add(tfk.layers.Conv2D(64, (3,3),
73                               activation=activation,
74                               kernel_initializer='he_uniform'))
75 net_TF.add(tfk.layers.MaxPool2D(pool_size=(2,2)))
76 net_TF.add(tfk.layers.Flatten())
77 net_TF.add(tfk.layers.Dense(2048, activation=activation,

```

```

78         kernel_initializer='he_uniform'))
79 net_TF.add(tfk.layers.Dense(10, activation='softmax',
80                             kernel_initializer='he_uniform'))
81
82 opt = tfk.optimizers.Adam(eta)
83 net_TF.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['
    accuracy'])
84
85 start = time()
86 """
87 net.train(x, y_train, batch_size=bs, epochs=ep)
88 """
89 t_train = time() - start
90
91 start = time()
92 net_TF.fit(x_TF, y_train, batch_size=bs, epochs=ep)
93 t_train_TF = time() - start
94
95 y_test = np.argmax(y_test, 1).T
96
97 y_tilde_TF = net_TF.predict(x_test.reshape(10000,28,28,1))
98 guess_TF = np.argmax(y_tilde_TF, 1).T
99 print('Accuracy with TensorFlow=', np.sum(guess_TF == y_test)/100)
100 print('Time needed for training:', t_train_TF)
101
102 # Our implementation: 3630s (60,5 min)
103 # Tensorflow: 206 s (3,4 min)

```

In my case, the network needed approx 3600s (60 min) for training 10 epochs, while TensorFlow only needed approx 200s (3,3 min). If you are able to run CUDA on your system, TensorFlow can be significantly faster.