
Mathematics of Neural Networks
winter semester 2021/2022
exercise sheet 8

Exercise 1: (4 points)

- a) Implement the computation of CAM.
- b) (Optional) Implement the computation of GradCAM.

Use the following skeleton and use the test as provided in the skeleton. Exchange `truck.jpg` to try different images.

```
1 import numpy as np
2 import tensorflow as tf
3 import matplotlib.pyplot as plt
4
5 from tensorflow import keras as K
6 from scipy.ndimage import zoom
7 from matplotlib import cm
8
9 preprocess_input = K.applications.resnet50.preprocess_input
10 decode_predictions = K.applications.resnet50.decode_predictions
11
12 # Load the ResNet50 network
13 def get_model():
14
15     return K.applications.resnet50.ResNet50()
16
17 """
18 Load an image from a given path with a given size
19 as TensorFlow Tensor.
20 """
21 def get_image(img_path, size):
22
23     # Load image
24     img = K.preprocessing.image.load_img(img_path, target_size=size)
25     # Cast img to a numpy array with shape (3, size[0], size[1])
26     img = K.preprocessing.image.img_to_array(img)
27     # Transform img to a 4-tensor of shape (1, 3, size[0], size[1])
28     img = tf.expand_dims(img, axis=0)
29     # Cast to float32, if not done yet
30     img = tf.cast(img, dtype=tf.float32)
31
32     return preprocess_input(img)
33
34 """
35 For a given model with last convolutional layer
36 last_conv_layer and given input x get the output
37 of the model and the last convolutional layer
38 """
```

```

39 def get_output(model, last_conv_layer, x):
40
41     get = K.Model([model.layers[0].input],
42                   [last_conv_layer.output,
43                    model.layers[-1].output])
44
45     return get(x)
46
47 # analyze the predictions of the net
48 def analyze_predictions(predictions):
49     predicted_class = predictions.argmax()
50     top5 = decode_predictions(predictions, top=5)[0]
51     print("=====\resulting\top\predictions:\=====")
52     for i in range(5):
53         print("{}:\probability\{:6.2f}%\for\the\class\{}\\"
54               .format(i + 1, 100 * top5[i][2], top5[i][1]))
55     return predicted_class
56
57 """
58 Get the CAM heatmap for a given model, where the name
59 of the last convolutional layer is last_conv_layer_name
60 """
61 def get_cam_heatmap(x, model, last_conv_layer_name):
62
63     last_conv_layer = model.get_layer(last_conv_layer_name)
64
65     # Get the output of the last convolutional layer and the
66     # predicted class
67     last_conv_out, predictions = get_output(model, last_conv_layer, x)
68     predicted_class = analyze_predictions(predictions.numpy())
69
70     # Get the weights of the prediction layer
71     W, b = model.layers[-1].get_weights()
72
73     """
74     TODO Sum the output of the last convolutional layer
75           over the channels. Scale each channel with the
76           corresponding weight of the predicted class
77           (See lecture notes, p.129 (version from 11.01.21)).
78     """
79     heatmap = None
80
81     # The returned heatmap should have shape (h,w), where h is the
82     # height of the output of the last convolutional layer, and w is
83     # its width.
84     """
85     TODO Assure that the heatmap is a 2D numpy array. You can convert
86           tensorflow tensors to numpy arrays with the numpy method.
87           If you want a numpy array from the tensor x, call c.numpy().
88           If you have a 4D array x of shape (1, h, w, 1), you can
89           replace it by x[0, :, :, 0] to obtain a 2D array.
90     """
91
92     return heatmap
93
94 def get_gradient_cam_heatmap(img, model, last_conv_layer_name,

```

```

95         classifier_layer_names):
96
97     last_conv_layer = model.get_layer(last_conv_layer_name)
98
99     # Splitting the network into convolutional and classifier part
100    # Model for the convolutional part
101    convolutional_model = K.Model(model.inputs, last_conv_layer.output)
102
103    # Model for the classifier part
104    classifier_input = K.Input(shape=last_conv_layer.output.shape[1:])
105    x = classifier_input
106    for layer_name in classifier_layer_names:
107        x = model.get_layer(layer_name)(x)
108
109    classifier_model = K.Model(classifier_input, x)
110
111    # Track gradients with GradientTape()
112    with tf.GradientTape() as tape:
113
114        # Call the convolutional part with img as input
115        last_conv_out = convolutional_model(img)
116
117        # Track the derivatives with respect to the output of the
118        # last convolutional layer.
119        tape.watch(last_conv_out)
120
121        # Call the classifier part with last_conv_out as input
122        predictions = classifier_model(last_conv_out)
123
124        # Analyze the predictions and get winning class data
125        predicted_class = analyze_predictions(predictions.numpy())
126        top_class_channel = predictions[:,predicted_class]
127
128        # Get the derivatives of the predicted class channel w.r.t. the output
129        # of the last convolutional layer
130        gradient = tape.gradient(top_class_channel, last_conv_out)
131
132        # gradient is of shape (1, oh, ow, c) where oh, ow are the height and
133        # width of the outout of the last convolutional layer. The average over
134        # the first three axes has to be taken.
135        pooled_gradient = tf.reduce_mean(gradient, axis=(0,1,2))
136
137        """
138        TODO Sum the output of the last convolutional layer over the channels.
139        Scale each channel with the corresponding pooled gradient
140        (see lecture notes, p.129 (version from 01.11.21)).
141        """
142        heatmap = None
143
144        # The returned heatmap should have shape (h,w), where h is the
145        # height of the output of the last convolutional layer, and w is
146        # its width.
147        """
148        TODO Assure that the heatmap is a 2D numpy array. You can convert
149        tensorflow tensors to numpy arrays with the numpy method.
150        If you want a numpy array from the tensor x, call c.numpy().

```

```

151         If you have an 4D array x of shape (1, h, w, 1), you can
152         replace it by x[0, :, :, 0] to obtain a 2D array.
153     """
154
155     return heatmap
156
157 def superimpose_heatmap(img_path, heatmap, alpha=.7):
158
159     # load image, e.g., float array of shape (465, 621, 3)
160     image_np = plt.imread(img_path).astype(np.float32)
161     heatmap_uint8 = np.uint8(np.maximum(heatmap, 0) / heatmap.max() * 255)
162     cm_jet = cm.get_cmap("jet")
163     jet_colors = cm_jet(np.arange(256))[:, :3]
164     heatmap_jet = jet_colors[heatmap_uint8]
165     # scale color heatmap to shape (465, 621, 3)
166     target_h, target_w, _ = image_np.shape
167     h, w, _ = heatmap_jet.shape
168     heatmap_scaled = zoom(heatmap_jet, (target_h/h, target_w/w, 1))
169     heatmap_scaled_uint8 = np.uint8(np.maximum(heatmap_scaled, 0)
170                                     / heatmap_scaled.max() * 255)
171     # superimpose image and heatmap
172     return np.uint8(image_np * (1 - alpha) + heatmap_scaled_uint8 * alpha)
173
174 # show image using matplotlib
175 def show_image(image_path, title=None):
176     image = plt.imread(image_path)
177     plt.imshow(image)
178     plt.title(title)
179     plt.show()
180
181 # show heatmap using matplotlib and cm
182 def show_heatmap(heatmap, title=None):
183     plt.imshow(heatmap, cmap='jet')
184     plt.title(title)
185     plt.show()
186
187
188 if __name__ == '__main__':
189
190     img_path = 'truck.jpg'
191     resnet_size = (224,224)
192
193     model = get_model()
194
195     last_conv_layer_name = 'conv5_block3_out'
196     classifier_layer_names = ['avg_pool',
197                               'predictions']
198
199     img = get_image(img_path, resnet_size)
200
201     heatmap_cam = get_cam_heatmap(img, model, last_conv_layer_name)
202     heatmap_grad_cam = get_gradient_cam_heatmap(img, model,
203                                                  last_conv_layer_name,
204                                                  classifier_layer_names)
205
206     image_cam = superimpose_heatmap(img_path, heatmap_cam)

```

```

207 image_grad_cam = superimpose_heatmap(img_path, heatmap_grad_cam)
208
209 plt.axis('off')
210 plt.imshow(image_cam)
211 plt.savefig('truck_cam.pdf')
212 plt.close()
213
214 plt.axis('off')
215 plt.imshow(image_grad_cam)
216 plt.savefig('truck_grad_cam.pdf')
217 plt.close()

```

Exercise 2: (2 points) Let \mathbf{B}_0 be symmetric. Let $\mathbf{y}_k, \mathbf{s}_k \in \mathbb{R}^n$ be given for all $k \in \mathbb{N}_0$. Prove that all BFGS approximations

$$\mathbf{B}_{k+1} = \mathbf{B}_k + \frac{\mathbf{y}_k \mathbf{y}_k^\top}{\mathbf{y}_k^\top \mathbf{s}_k} - \frac{\mathbf{B}_k \mathbf{s}_k \mathbf{s}_k^\top \mathbf{B}_k^\top}{\mathbf{s}_k^\top \mathbf{B}_k \mathbf{s}_k}, \quad k \in \mathbb{N}_0$$

to the Hessian are symmetric and satisfy the secant condition $\mathbf{B}_{k+1} \mathbf{s}_k = \mathbf{y}_k$.

Exercise 3: (2+2*+2 points)

- a) Let regular $\mathbf{A} \in \mathbb{R}^{n \times n}$ and $\mathbf{U} \in \mathbb{R}^{n \times k}$, $\mathbf{V} \in \mathbb{R}^{k \times n}$ be given. Prove the Sherman-Morrison-Woodbury formula

$$(\mathbf{A} + \mathbf{UV})^{-1} = \mathbf{A}^{-1} - \mathbf{A}^{-1} \mathbf{U} (\mathbf{I}_k + \mathbf{VA}^{-1} \mathbf{U})^{-1} \mathbf{VA}^{-1}$$

for regular $\mathbf{I}_k + \mathbf{VA}^{-1} \mathbf{U}$.

- b) Use the Sherman-Morrison-Woodbury formula to derive the update

$$\mathbf{H}_{k+1} = \left(\mathbf{I} - \frac{\mathbf{s}_k \mathbf{y}_k^\top}{\mathbf{y}_k^\top \mathbf{s}_k} \right) \mathbf{H}_k \left(\mathbf{I} - \frac{\mathbf{y}_k \mathbf{s}_k^\top}{\mathbf{y}_k^\top \mathbf{s}_k} \right) + \frac{\mathbf{s}_k \mathbf{s}_k^\top}{\mathbf{y}_k^\top \mathbf{s}_k}$$

of the inverse $\mathbf{H}_k = \mathbf{B}_k^{-1}$ of the approximation \mathbf{B}_k of the Hessian in BFGS from [Exercise 2](#).

- c) Suppose that the vectors $\mathbf{s}_k, \mathbf{y}_k$ of [Exercise 2](#) satisfy additionally the curvature condition $\mathbf{s}_k^\top \mathbf{y}_k > 0$ and that \mathbf{B}_0 is symmetric positive definite. Prove that then all \mathbf{H}_k (and thus all \mathbf{B}_k) are symmetric positive definite¹.

Exercise 4: (3 points) Implement Newton's method, BFGS, and L-BFGS for the minimization of the function

$$f: \mathbb{R}^n \rightarrow \mathbb{R}, \quad f(\mathbf{x}) = \frac{1}{2} \left((x_1 - 1)^2 + \sum_{k=2}^n k(x_k - x_{k-1})^2 \right)$$

How well do these methods work? For L-BFGS compute the update with the following algorithm. Use $\gamma \in \{1, \mathbf{s}_{k-1}^\top \mathbf{y}_{k-1} / \mathbf{y}_{k-1}^\top \mathbf{y}_{k-1}\}$.

¹A matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ is symmetric positive definite when it is symmetric and for all nonzero vectors $\mathbf{v} \in \mathbb{R}^n$ we have $\mathbf{v}^\top \mathbf{A} \mathbf{v} > 0$.

Algorithm 1 Iterative computation of the update in L-BFGS

```

1:  $\mathbf{p} \leftarrow -\mathbf{g}_k$ ;
2: for  $j = k - 1 : k - m : -1$  do
3:    $\alpha_j \leftarrow \rho_j \mathbf{s}_j^\top \mathbf{p}$ ;
4:    $\mathbf{p} \leftarrow \mathbf{p} - \mathbf{y}_j \alpha_j$ ;
5: end for
6:  $\mathbf{p} \leftarrow \mathbf{p} \gamma$ ;
7: for  $j = k - m : k$  do
8:    $\alpha_j \leftarrow \alpha_j - \rho_j \mathbf{y}_j^\top \mathbf{p}$ ;
9:    $\mathbf{p} \leftarrow \mathbf{p} + \mathbf{s}_j \alpha_j$ ;
10: end for

```

Hint: Since f is quadratic, the gradient $\nabla f(\mathbf{x})$ is linear in each component. The Hessian $\nabla^2 f(\mathbf{x})$ is constant, symmetric, and positive definite.